

- **Definition**
- **Key Interview Points**
- **Mnemonic**
- **Story**
- **Interview Tip**
- **Example/Model File**

Docker Compose (`docker-compose.yml`)

Definition Docker Compose is a tool for defining and running multi-container Docker applications via a single YAML file that configures services, networks, and volumes ([Docker Documentation][1]).

Key Interview Points

- The top-level `version` specifies the Compose file format (v2, v3, or the Compose Specification) ([Docker Documentation][2]).
- Under `services`, each key is a service name; its value defines either an `image` or a `build` context, along with `ports`, `environment`, and `depends_on` ([Docker Documentation][3]).
- `networks` allow services to communicate on custom bridges; `volumes` define persistent storage shared across services ([Docker Documentation][4]).

Mnemonic

"Very Silly Penguins Wiggle"

- *Version*
- *Services*
- *Ports/Parameters*
- *Networks*
- *Volumes*

Story Imagine a penguin colony throwing an ice-party: they pick the **version** of their music playlist, set up **services** (hot cocoa stand, ice slide), map **ports** (ice tunnels) and **parameters** (guest list), string **networks** of glowing lights, and stash leftovers in giant **volumes** of ice.

Interview Tip

"In `docker-compose.yml`, I always begin with `version`, define each service under `services`, map its `ports` and `environment` variables, then configure any custom `networks`, and finally declare `volumes` for data persistence."

Example (`docker-compose.yml`)

```
version: '3.8'

services:
  web:
    image: nginx:stable
    ports:
      - "80:80"
    environment:
      - NGINX_HOST=example.com
      - NGINX_PORT=80
```

```

api:
  build:
    context: ./api
    dockerfile: Dockerfile
  ports:
    - "4000:4000"
  depends_on:
    - db

db:
  image: postgres:13
  restart: always
  environment:
    POSTGRES_USER: user
    POSTGRES_PASSWORD: pass
  volumes:
    - db-data:/var/lib/postgresql/data

networks:
  default:
    driver: bridge

volumes:
  db-data:

```

GitHub Actions Workflow (`.github/workflows/ci.yml`)

Definition GitHub Actions workflows automate CI/CD by running jobs in response to repository events, defined in YAML under `.github/workflows/` ([GitHub Docs][5]).

Key Interview Points

- `name`: a human-readable identifier for the workflow.
- `on`: specifies triggers such as `push`, `pull_request`, or a cron schedule.
- `jobs`: each job runs in isolation; uses `runs-on` to pick a runner (e.g., `ubuntu-latest`), and contains ordered `steps` (actions or shell commands) ([GitHub Docs][6]).

Mnemonic

"Naughty Otters Jump Right Swiftly"

- *Name*
- *On*
- *Jobs*
- *Runs-on*
- *Steps*

Story Picture a troupe of mischievous otters putting on a river show: they give the performance a **name**, decide **on** which days to perform, assign **jobs** like juggling and diving, choose the rafts they'll **run-on**, and rehearse the **steps** of their synchronized splash finale.

Interview Tip

"In my GitHub Actions YAML, I define `name`, set the `on` triggers, declare `jobs` with specific `runs-on` environments, and list each `steps` block—often starting with `checkout`, `setup`, `build`, `test`, and `deploy` actions."

Example (ci.yml)

```
name: CI Pipeline

on:
  push:
    branches: [ main ]
  pull_request:
    branches: [ main ]

jobs:
  build-and-deploy:
    runs-on: ubuntu-latest

    steps:
      - name: Checkout code
        uses: actions/checkout@v3

      - name: Set up Node.js
        uses: actions/setup-node@v3
        with:
          node-version: '18'

      - name: Install dependencies
        run: npm ci

      - name: Run tests
        run: npm test

      - name: Build Docker image
        run: docker build -t myapp:${{ github.sha }} .

      - name: Push to AWS ECR
        env:
          AWS_REGION: us-east-1
        run: |
          aws ecr get-login-password --region $AWS_REGION \
          | docker login --username AWS --password-stdin
          123456789012.dkr.ecr.$AWS_REGION.amazonaws.com
          docker tag myapp:${{ github.sha }}
          123456789012.dkr.ecr.$AWS_REGION.amazonaws.com/myapp:latest
          docker push 123456789012.dkr.ecr.$AWS_REGION.amazonaws.com/myapp:latest
```

Jenkins Pipeline (Jenkinsfile)

Definition Pipeline-as-code for Jenkins, defining your build/test/deploy workflow in a Groovy DSL.

Key Interview Points

- Top-level `pipeline {}` block encloses everything.
- `agent` declares where the pipeline runs (e.g. `any`, `label 'docker'`).
- `stages` groups named `stage('Build') { steps { ... } }`, etc.
- Within each `steps`, you run shell commands or use plugins.
- You can parallelize stages and manage credentials via `environment` and `credentials`.

Mnemonic

"Purple Ants Stroll Softly"

- *Pipeline* → *Agent* → *Stages* → *Stage* → *Steps*

Story Purple ants form a construction crew: they lay the pipeline trunk, each ant (agent) grabs materials, map out stages of the build, work through each stage's steps, then celebrate when the bridge is complete.

Interview Tip

"My Jenkinsfile begins with `pipeline { agent any }`, defines stages for Build, Test, and Deploy, and uses shared library steps for environment setup and artifact archiving."

Example (Jenkinsfile)

```
pipeline {
  agent any
  environment {
    REGISTRY = '123456789012.dkr.ecr.us-east-1.amazonaws.com'
    IMAGE    = "${REGISTRY}/myapp"
  }
  stages {
    stage('Checkout') {
      steps { checkout scm }
    }
    stage('Build & Test') {
      steps {
        sh 'npm ci'
        sh 'npm test'
      }
    }
    stage('Docker Build & Push') {
      steps {
        sh """
            $(aws ecr get-login --no-include-email --region us-east-1)
            docker build -t $IMAGE:${env.BUILD_NUMBER} .
            docker push $IMAGE:${env.BUILD_NUMBER}
          """
      }
    }
    stage('Deploy') {
      steps {
        sh 'kubectl apply -f k8s/deployment.yaml'
```

```
    sh 'kubectl rollout status deployment/myapp'
  }
}
}
```

Terraform Module (main.tf)

Definition Terraform is a declarative IaC tool that provisions cloud infrastructure via `.tf` files and a state backend.

Key Interview Points

- `terraform init` bootstraps providers.
- `terraform plan` previews changes.
- `terraform apply` executes.
- `terraform destroy` tears down resources.
- Use remote state (S3 + DynamoDB) for team collaboration.

Mnemonic

“Pandas Really Value Oranges”

- ***Provider → Resource → Variable → Output***

Story Pandas run a jungle café: they choose their bamboo crate **provider**, list each fruit **resource**, set pricing **variables**, and update the earnings **outputs** on their chalkboard.

Interview Tip

“In Terraform I start with a `provider` block, declare `resource` blocks for each AWS component, parameterize settings with `variable`, and expose useful IDs via `output`.”

Example (main.tf)

```
terraform {
  required_providers {
    aws = { source = "hashicorp/aws", version = ">= 4.0" }
  }
  backend "s3" {
    bucket = "my-terraform-state"
    key    = "eks/terraform.tfstate"
    region = "us-east-1"
  }
}

provider "aws" {
  region = "us-east-1"
}

variable "app_name" {
  type    = string
  default = "myapp"
```

```

}

resource "aws_vpc" "main" {
  cidr_block = "10.0.0.0/16"
  tags      = { Name = "${var.app_name}-vpc" }
}

output "vpc_id" {
  value = aws_vpc.main.id
}

```

Ansible Playbook (site.yml)

Definition YAML-based automation for configuration management and app deployment, agentless over SSH.

Key Interview Points

- Top-level list of plays with `hosts` .
- `become: yes` for privilege escalation.
- `vars` for reusable parameters.
- `tasks` list actions (modules).
- `handlers` run on notification (e.g., service restarts).

Mnemonic

"Hungry Bears Value Tasty Honey"

- *Hosts* → *Become* → *Vars* → *Tasks* → *Handlers*

Story Hungry bears pick their berry bush **hosts**, choose to **become** stronger, measure berry **vars**, perform picking **tasks**, and call **handlers** (the bees) if the honey jars break.

Interview Tip

"My playbooks start with `hosts: all` , optionally `become: yes` , define `vars` , list `tasks` like package installs and file templates, and include `handlers` for service reloads."

Example (site.yml)

```

- hosts: webservers
  become: yes
  vars:
    app_repo: https://github.com/example/myapp.git
    app_dest: /var/www/myapp

  tasks:
    - name: Install NGINX
      apt:
        name: nginx
        state: latest
        update_cache: yes

```

```

- name: Checkout application code
  git:
    repo: "{{ app_repo }}"
    dest: "{{ app_dest }}"
    version: main

- name: Configure NGINX
  template:
    src: nginx.conf.j2
    dest: /etc/nginx/sites-available/myapp
  notify:
    - reload nginx

handlers:
- name: reload nginx
  service:
    name: nginx
    state: reloaded

```

Kubernetes Deployment (deployment.yaml)

Definition A Deployment resource manages stateless pods with rolling updates and self-healing.

Key Interview Points

- `apiVersion: apps/v1` & `kind: Deployment` .
- `metadata.name` and labels.
- `spec.replicas` , `selector.matchLabels` .
- `template` block contains pod spec (containers, ports, env).
- Enables rolling updates and rollbacks out of the box.

Mnemonic

"Ancient Kings Make Simple Temples"

- *API Version → Kind → Metadata → Spec → Template*

Story Ancient kings declare the **apiVersion**, choose the **kind** of structure, engrave **metadata**, outline the **spec**, and lay the **template** foundation blocks to ensure the temple always remains intact.

Interview Tip

"In my Deployment YAML I use `apiVersion` and `kind`, set `metadata` labels, define the desired `replicas` and `selector`, and under `template` configure containers, ports, and liveness probes."

Example (deployment.yaml)

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp
  labels:

```

```

  app: myapp
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp
          image: 123456789012.dkr.ecr.us-east-1.amazonaws.com/myapp:latest
          ports:
            - containerPort: 80
          env:
            - name: NODE_ENV
              value: production

```

Helm Chart (Chart.yaml & values.yaml)

Definition Helm is the Kubernetes package manager; charts bundle multiple manifests and parameterize them with values.

Key Interview Points

- Chart.yaml contains chart metadata: apiVersion , name , version , dependencies .
- values.yaml holds default configuration that templates consume.
- Commands: helm repo add , helm install , helm upgrade , helm rollback .

Mnemonic

"All Noble Vampires Drink Vanilla"

- **API Version** → **Name** → **Version** → **Dependencies** → **Values**

Story Noble vampires host a midnight ball: they note the **apiVersion** of their society, announce their **name**, update fang **version**, invite **dependencies** (bat entourage), and sip **vanilla** elixir as they dance.

Interview Tip

"In Chart.yaml I set apiVersion , name , version , and any dependencies . I use values.yaml to override container image tags, replica counts, and environment-specific configs."

Example (Chart.yaml)

```

apiVersion: v2
name: myapp
description: A Helm chart for myapp
type: application
version: 0.1.0

```



```
appVersion: "1.0.0"
dependencies:
  - name: redis
    version: 15.0.0
    repository: https://charts.bitnami.com/bitnami
```

Example (values.yaml)

```
replicaCount: 2

image:
  repository: 123456789012.dkr.ecr.us-east-1.amazonaws.com/myapp
  tag: latest

service:
  type: LoadBalancer
  port: 80

resources: {}
nodeSelector: {}
tolerations: []
affinity: {}
```

ArgoCD Application (application.yaml)

Definition An ArgoCD Application declares a Git repo path to sync into one or more Kubernetes clusters, enabling GitOps.

Key Interview Points

- `apiVersion: argoproj.io/v1alpha1` & `kind: Application` .
- `metadata.name` , `metadata.namespace` .
- Under `spec : project` , `source` (`repoURL` , `targetRevision` , `path`) , `destination` (`server` , `namespace`) .
- `syncPolicy` can be `automated` (with `prune` and `selfHeal`) or `manual` .

Mnemonic

"Always Keep Managing Syncs"

- *API Version → Kind → Metadata → Spec → SyncPolicy*

Story ArgoCD is like a dream chef: you write your recipe (Git), it notes the **apiVersion**, selects the **kind** of dish, labels it in **metadata**, follows the **spec** for ingredients, and uses **syncPolicy** to auto-reheat whenever the menu changes.

Interview Tip

"An ArgoCD Application YAML starts with `apiVersion` and `kind` , includes `metadata` , then under `spec` defines the Git `source` , `destination` cluster, and a `syncPolicy` for automated reconciliation."

Example (application.yaml)

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: myapp
  namespace: argocd
spec:
  project: default
  source:
    repoURL: https://github.com/example/myapp
    targetRevision: main
    path: helm-chart
  destination:
    server: https://kubernetes.default.svc
    namespace: production
  syncPolicy:
    automated:
      prune: true
      selfHeal: true
```

Prometheus Config (prometheus.yml)

Definition Prometheus configuration defines how metrics are scraped, which alerting rules to load, and where to send long-term data.

Key Interview Points

- `global` : defaults such as `scrape_interval` and `evaluation_interval` .
- `scrape_configs` : list of jobs, each with `job_name` and service discovery or static targets.
- `rule_files` : paths to recording and alerting rules.
- `alerting` : configuration for Alertmanager endpoints.
- `remote_write` : send data to long-term storage.

Mnemonic

"Giant Sharks Race Around Reefs"

- ***Global*** → ***Scrape_configs*** → ***Rule_files*** → ***Alerting*** → ***Remote_write***

Story Giant sharks patrol the reef: they set **global** hunting grounds, follow **scrape_configs** for prey locations, memorize **rule_files** for safe zones, fire **alerting** signals when low on energy, and **remote_write** their catch logs to the deep-sea archives.

Interview Tip

"In prometheus.yml, I configure the global scrape interval, define scrape_configs for each target, reference rule_files for alerts, set up the alerting section to point to Alertmanager, and use remote_write for long-term storage in Thanos or Cortex."

Example (prometheus.yml)

```
global:
  scrape_interval: 15s
  evaluation_interval: 30s

scrape_configs:
  - job_name: 'kubernetes-nodes'
    kubernetes_sd_configs:
      - role: node

  - job_name: 'kubernetes-pods'
    kubernetes_sd_configs:
      - role: pod

rule_files:
  - 'alert.rules.yml'

alerting:
  alertmanagers:
    - static_configs:
      - targets:
        - 'alertmanager:9093'

remote_write:
  - url: 'https://long-term-storage.example.com/api/v1/write'
```